Bounded follow-up results

G2 minus G1 accuracy: +0.000000, paired conditional 95% interval [+0.000000, +0.000000]; achieved calls 2.186 versus 2.180.
G2 minus G1 calls: +0.006400, paired conditional 95% interval [-0.002405, +0.015800]; achieved calls 2.186 versus 2.180.
G2 minus G1 utility: -0.000064, paired conditional 95% interval [-0.000158, +0.000024]; achieved calls 2.186 versus 2.180.

The observed mean utility difference does not favor G2 under the predeclared 0.01 preference. This negative/null finding is retained without further redesign.

The final pool contains 1000 newly held-out official-training questions. Policy selection used only prior search/validation data. Evolutionary results average five seeds; whole-question intervals condition on fitted policies. All probes are budgeted. Mechanism diagnostics and simple ordered-triple/hill-climbing controls are retained regardless of outcome. This is a post-study repair within Study 2, not a new research direction or an originally preregistered study.